



UPPSALA UNIVERSITET

Assignment 3

Quicksort

Yuanjing Yang

June 26, 2025

1 Introduction and Problem Description

The goal of this assignment is to implement a parallel version of the Quicksort algorithm for MPI. We need to build several functions to read unsorted lists of integers and distribute them among different processes, doing *local_sort* function and *global_sort* function in sequences, gathering into the root processes, and printing it out. Here we evaluate different pivot selection strategies to analyze their impact on performance (strong scalability and weak scalability).

2 Algorithm and implementation

Table 1: Parallel Quicksort Algorithm Overview

Step	Description
1	The root process (rank 0) uses <code>read_input</code> to read the unsorted list of integers. The array is then divided into chunks and distributed to all processes using <code>MPI.Scatterv</code> .
2	Each process performs local sorting using the standard C function <code>qsort</code> .
3	Pivot selection strategies: <ul style="list-style-type: none">• <code>MEDIAN_ROOT</code>: Median of root process's local data.• <code>MEAN_MEDIAN</code>: Mean of local medians from all processes.• <code>MEDIAN_MEDIAN</code>: Median of the medians collected from all processes.
4	Each process splits its data into two parts: values smaller or larger than the pivot.
5	Partner processes exchange relevant data using <code>MPI.Sendrecv</code> , ensuring elements less than or equal to the pivot are sent to the lower group, and the rest to the upper group.
6	After recursion finishes, sorted segments are gathered on the root process using <code>MPI.Gatherv</code> .

3 Performance Experiment

I ran strong and weak scalability tests on UPPMAX using Rackham.

For strong scalability, I used fixed input size of 1 billion integers (input1000000000.txt), using processes 1, 2, 4, 8, 16 in 3 different strategies.

From the results below, we can see all three pivot strategies exhibit good strong scalability up to 16 processes. Efficiency remains relatively high (above 60%) even with 16 processes. Strategy 2 (`MEAN_MEDIAN`) and Strategy 3 (`MEDIAN_MEDIAN`) provide slightly better efficiency and speedup than Strategy

1, especially as the number of processes increases.

Table 2: Strong scalability for Strategy 1 (Median_ROOT)

Processes	Runtime	Speedup	Efficiency
1	237.122	100%	100%
2	121.984	194%	97%
4	67.2208	353%	88%
8	37.1413	634%	79%
16	25.8171	918%	57%

Table 3: Strong scalability for Strategy 2 (MEAN_MEDIAN)

Processes	Runtime	Speedup	Efficiency
1	238.96	100%	100%
2	123.5	193%	97%
4	66.277	361%	90%
8	38.2289	625%	78%
16	24.0662	993%	62%

Table 4: Strong scalability for Strategy 3 (MEDIAN_MEDIAN)

Processes	Runtime	Speedup	Efficiency
1	237.144	100%	100%
2	122.332	194%	97%
4	65.9955	362%	90%
8	37.3793	634%	79%
16	23.8909	993%	62%




Figure 1: The Speedup of strong scalability




Figure 2: The Efficiency of strong scalability

For weak scalability, I used input scaled with the number of process(125 million integers per process), using processes 1, 2, 4, 8, 16 in 3 different strategies. As the number of processes increases, the efficiency drops significantly. This is expected due to increased communication overhead and imbalance during the recursive splitting and merging phases. All three pivot strategies show similar weak scalability behavior.

Table 5: Weak scalability for Strategy 1 (Median_ROOT)

Processes	Runtime	Speedup	Efficiency
1	46.083	100%	100%
2	67.986	93%	46%
4	110.462	83%	21%
8	198.888	70%	9%
16	400.956	53%	3%

Table 6: Weak scalability for Strategy 2 (MEAN_MEDIAN)

Processes	Runtime	Speedup	Efficiency
1	26.305	100%	100%
2	28.7194	93%	46%
4	31.2929	84%	21%
8	38.131	71%	9%
16	48.7517	54%	3%

Table 7: Weak scalability for Strategy 3 (MEDIAN_MEDIAN)

Processes	Runtime	Speedup	Efficiency
1	26.305	100%	100%
2	28.7194	92%	46%
4	31.2929	84%	21%
8	38.131	69%	9%
16	48.7517	54%	3%




Figure 3: The Runtime of weak scalability




Figure 4: The Speedup of weak scalability

While the speedup values are comparable across strategies, Strategy 2 and 3 maintain better efficiency under strong scalability scenarios. For weak scalability, the difference is minimal. Strategy 1 (Median in Root) is the simplest but results in less balanced partitions and worse scaling at higher process counts.

4 Discussion

The parallel quicksort algorithm implemented in this assignment shows the benefits and drawback of parallelization when applied to recursive sorting.

For strong scalability, the results indicate good strong scalability all three pivot strategies, we can see the execution time decrease as we expected. With 8 processes, all strategies achieve more than 6 times, and the efficiency remains high below 8 processes. This result comes from the increasing communication overhead and load imbalance during recursive process splitting, and strategy 1 is worse than other strategies, it can be reasonable due to it is the simplest and direct strategy, this suggests that balanced pivot selection is important for good scalability.

For weak scalability, as the number of processes and the problem size grow, execution time increases, which also can be attributed to communication overhead and frequent memory allocation and data merging, but the performance of strategy 2 and 3 still better than strategy 1, due to producing more balanced partitions and reducing the depth of recursion.

5 Conclusion

From the result of performance in strong scalability and weak scalability, we can see that a more balanced pivot strategy such as *MEAN_MEDIAN* and *MEDIAN_MEDIAN* show good strong scalability, and when we tried it on *fixed_size* problems (weak scalability), it does not scale efficiently, the main factors are recursive communication overhead, even *MEDIAN_ROOT* can lead to minimal communication overhead, but in my experiment, balanced partitioning resulted in fewer recursion levels and more evenly distributed workload, the additional communication overhead was offset by the improved overall parallel efficiency.

6 Appendix

```
1 quicksort.c:
2 #include "quicksort.h"
3 #include "pivot.h"
4 #include <stdlib.h>
5 #include <string.h>
6 #include <stdio.h>
7 #include <mpi.h>
8
9
10 #define NOPRINTING
11
12 int check_and_print(int *elements, int n, char *file_name){
13     int sort_element=sorted_ascending(elements, n);
```

```

14     if(!sort_element){
15         printf("Error: the elements are not sorted in
           ascending order.\n");
16     }
17     FILE *file=fopen(file_name, "w");
18     if (!file) return -1;
19
20     for(int i=0; i<n; i++){
21         fprintf(file, "%d", elements[i]);
22         if(i < n-1) fprintf(file, " ");
23     }
24     fprintf(file, "\n");
25     fclose(file);
26     return 0;
27 }
28
29 int distribute_from_root(int *all_elements, int n, int **
    local_elements){
30     int rank, size;
31     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
32     MPI_Comm_size(MPI_COMM_WORLD, &size);
33
34     int base_elements=n/size;
35     int remainder=n%size;
36     int *counts=malloc(size*sizeof(int));
37     int *chunk_start=malloc(size*sizeof(int));
38
39     int offset=0;
40     for(int i=0; i<size; i++){
41         counts[i]=base_elements+(i<remainder? 1:0);
42         chunk_start[i]=offset;
43         offset+=counts[i];
44     }
45
46     int local_n=counts[rank];
47     *local_elements=malloc(local_n*sizeof(int));
48     MPI_Scatterv(all_elements, counts, chunk_start, MPI_INT,
         *local_elements, local_n, MPI_INT, 0, MPI_COMM_WORLD
         );
49
50     free(counts);
51     free(chunk_start);
52
53     return local_n;
54 }
55
56 void gather_on_root(int *all_elements, int *local_elements,
    int local_n) {
57     int rank, size;
58     MPI_Comm_rank(MPI_COMM_WORLD, &rank);

```



```

59 MPI_Comm_size(MPI_COMM_WORLD, &size);
60
61 int *counts = NULL;
62 int *displs = NULL;
63
64 if (rank == 0) {
65     counts = malloc(size * sizeof(int));
66     displs = malloc(size * sizeof(int));
67 }
68
69 MPI_Gather(&local_n, 1, MPI_INT, counts, 1, MPI_INT, 0,
70           MPI_COMM_WORLD);
71
72 if (rank == 0) {
73     displs[0] = 0;
74     for (int i = 1; i < size; i++) {
75         displs[i] = displs[i-1] + counts[i-1];
76     }
77
78 MPI_Gatherv(local_elements, local_n, MPI_INT,
79             all_elements, counts, displs, MPI_INT,
80             0, MPI_COMM_WORLD);
81
82 if (rank == 0) {
83     free(counts);
84     free(displs);
85 }
86 }
87
88 int global_sort(int **elements, int n, MPI_Comm comm, int
89                pivot_strategy) {
90     int rank, size;
91     MPI_Comm_rank(comm, &rank);
92     MPI_Comm_size(comm, &size);
93
94     if (size == 1) {
95         return n;
96     }
97
98     // We need to check if size is even before proceeding
99     if (size % 2 != 0) {
100         if (rank == 0) {
101             fprintf(stderr, "Error: Number of processes (%d)
102                        must be even at all recursion levels.\n",
103                        size);
104         }
105         MPI_Abort(MPI_COMM_WORLD, 1);
106     }

```

```

105     int actual_pivot_value;
106     // We need the index of the first element greater than
107     pivot_value
108     int pivot_split_idx = select_pivot(pivot_strategy, *
        elements, n, comm, &actual_pivot_value);
109     int half_size = size / 2;
110     int new_color;
111     int partner_rank;
112     // here I spplit the data into send and keep two groups
113     int send_n;
114     int keep_n;
115     int *send_ptr;
116     int *keep_ptr;
117
118     if (rank < half_size) {
119         new_color = 0;
120         partner_rank = rank + half_size;
121
122         // Here the Elements <= pivot_value are kept,
            elements > pivot_value are sent
123         keep_n = pivot_split_idx;
124         send_n = n - pivot_split_idx;
125         keep_ptr = *elements;
126         send_ptr = *elements + pivot_split_idx;
127
128     } else {
129         new_color = 1;
130         partner_rank = rank - half_size;
131
132         // Elements > pivot_value are kept, elements <=
            pivot_value are sent
133         keep_n = n - pivot_split_idx;
134         send_n = pivot_split_idx;
135         keep_ptr = *elements + pivot_split_idx;
136         send_ptr = *elements;
137     }
138
139     // Here I exchange sizes first
140     int recv_n;
141     MPI_Sendrecv(&send_n, 1, MPI_INT, partner_rank, 0,
142         &recv_n, 1, MPI_INT, partner_rank, 0,
143         comm, MPI_STATUS_IGNORE);
144
145     // Allocating receive buffer
146     int* received_elements = (int*)malloc((recv_n > 0 ?
        recv_n : 1) * sizeof(int));
147     if (recv_n > 0 && !received_elements) {
148         fprintf(stderr, "Rank_%d: Malloc for
            received_elements failed\n", rank);

```

```

149     MPI_Abort(MPI_COMM_WORLD, 1);
150 }
151
152 // Here I exchange actual data
153 MPI_Sendrecv(send_ptr, send_n, MPI_INT, partner_rank, 1,
154             received_elements, recv_n, MPI_INT,
155             partner_rank, 1,
156             comm, MPI_STATUS_IGNORE);
157
158 // Begin merge
159 int new_n = keep_n + recv_n;
160 int *merged_elements = (int*)malloc((new_n > 0 ? new_n :
161                                     1) * sizeof(int));
162 if (new_n > 0 && !merged_elements) {
163     fprintf(stderr, "Rank%d: Malloc for merged_elements
164                failed\n", rank);
165     MPI_Abort(MPI_COMM_WORLD, 1);
166 }
167 merge_ascending(keep_ptr, keep_n, received_elements,
168                recv_n, merged_elements);
169
170 if (received_elements) free(received_elements);
171
172 free(*elements);
173 *elements = merged_elements;
174
175 // Splitting communicator
176 MPI_Comm new_comm;
177 MPI_Comm_split(comm, new_color, rank, &new_comm);
178 // Recursive call
179 int result_n = global_sort(elements, new_n, new_comm,
180                             pivot_strategy);
181 MPI_Comm_free(&new_comm);
182 return result_n;
183 }
184
185 void merge_ascending(int *v1, int n1, int *v2, int n2, int *
186 result){
187     int i = 0, j = 0, k = 0;
188     while (i < n1 && j < n2) {
189         if (v1[i] <= v2[j]) {
190             result[k++] = v1[i++];
191         } else {
192             result[k++] = v2[j++];
193         }
194     }
195     while (i < n1) {
196         result[k++] = v1[i++];
197     }
198     while (j < n2) {

```

```

193         result[k++] = v2[j++];
194     }
195 }
196
197 int read_input(char *file_name, int **elements) {
198     FILE *file = fopen(file_name, "r");
199     if (!file) {
200         perror("Couldn't open input file");
201         return -1;
202     }
203     int num_values;
204     if (fscanf(file, "%d", &num_values) != 1) {
205         perror("Couldn't read element count from input file");
206         fclose(file);
207         return -1;
208     }
209     *elements = malloc(num_values * sizeof(int));
210
211     if (!(*elements) && num_values > 0) {
212         perror("Memory allocation failed");
213         fclose(file);
214         return -1;
215     }
216
217     for (int i = 0; i < num_values; i++) {
218         if (fscanf(file, "%d", &((*elements)[i])) != 1) {
219             perror("Couldn't read elements from input file");
220             ;
221             free(*elements);
222             *elements = NULL;
223             fclose(file);
224             return -1;
225         }
226     }
227     fclose(file);
228     return num_values;
229 }
230
231 int sorted_ascending(int *elements, int n) {
232     for (int i = 1; i < n; i++) {
233         if (elements[i] < elements[i-1]) {
234             printf("Error at index %d: %d > %d\n", i - 1,
235                 elements[i - 1], elements[i]);
236             return 0;
237         }
238     }
239     return 1;

```

```

240
241
242 void swap(int *e1, int *e2) {
243     int tmp = *e1;
244     *e1 = *e2;
245     *e2 = tmp;
246 }
247
248 int main(int argc, char* argv[]) {
249     MPI_Init(&argc, &argv);
250     int rank, size;
251     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
252     MPI_Comm_size(MPI_COMM_WORLD, &size);
253
254     if (argc != 4) {
255         if (rank == 0)
256             printf("Usage: %s <input_file> <output_file> <
                pivot_strategy>\n", argv[0]);
257         MPI_Finalize();
258         return 1;
259     }
260
261     char *input_name = argv[1];
262     char *output_name = argv[2];
263     int pivot_strategy = atoi(argv[3]);
264
265     int* all_elements = NULL;
266     int* local_elements = NULL;
267     int total_n = 0;
268
269     double overall_start_time = MPI_Wtime();
270
271     if (rank == 0) {
272         total_n = read_input(input_name, &all_elements);
273         if (total_n <= 0) {
274             if (all_elements) free(all_elements);
275             MPI_Abort(MPI_COMM_WORLD, 1);
276         }
277     }
278
279     MPI_Bcast(&total_n, 1, MPI_INT, 0, MPI_COMM_WORLD);
280
281     if (total_n <= 0) {
282         MPI_Finalize();
283         return 1;
284     }
285
286     double distrubution_start_time = MPI_Wtime();
287     int local_n = distribute_from_root(all_elements, total_n
        , &local_elements);

```

```

288     double distrubution_end_time = MPI_Wtime();
289     double current_distr_time = distrubution_end_time -
        distrubution_start_time;
290     double local_serial_sort_start_time = MPI_Wtime();
291     if (local_n > 1) {
292         qsort(local_elements, local_n, sizeof(int), compare)
        ; //         compare
293     }
294     double local_serial_sort_end_time = MPI_Wtime();
295     double current_process_serial_time =
        local_serial_sort_end_time -
        local_serial_sort_start_time;
296     MPI_Barrier ( MPI_COMM_WORLD );
297     double global_sort_start_time = MPI_Wtime();
298     int sorted_n = global_sort(&local_elements, local_n,
        MPI_COMM_WORLD, pivot_strategy);
299     double global_sort_end_time = MPI_Wtime();
300     double current_process_global_sort_time =
        global_sort_end_time - global_sort_start_time;
301     MPI_Barrier ( MPI_COMM_WORLD );
302     if (rank == 0) {
303         free(all_elements);
304     }
305     all_elements = NULL;
306     if (rank == 0) {
307         all_elements = malloc(total_n * sizeof(int));
308     }
309     double gather_start_time=MPI_Wtime();
310     gather_on_root(all_elements, local_elements, sorted_n);
311     double gather_end_time=MPI_Wtime();
312     double current_process_gather_time = gather_end_time -
        gather_start_time;
313
314     double overall_end_time = MPI_Wtime();
315     double current_process_overall_time = overall_end_time -
        overall_start_time;
316
317     if (local_elements) free(local_elements);
318
319     double max_distr_time;
320     MPI_Reduce(&current_distr_time, &max_distr_time, 1,
        MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
321
322
323     double max_gather_time;
324     MPI_Reduce(&current_process_gather_time, &
        max_gather_time, 1, MPI_DOUBLE, MPI_MAX, 0,
        MPI_COMM_WORLD);
325
326

```

```

327     double max_overall_time;
328     MPI_Reduce(&current_process_overall_time, &
               max_overall_time, 1, MPI_DOUBLE, MPI_MAX, 0,
               MPI_COMM_WORLD);
329
330     double max_serial_time;
331     MPI_Reduce(&current_process_serial_time, &
               max_serial_time, 1, MPI_DOUBLE, MPI_MAX, 0,
               MPI_COMM_WORLD);
332
333     double max_global_sort_time;
334     MPI_Reduce(&current_process_global_sort_time, &
               max_global_sort_time, 1, MPI_DOUBLE, MPI_MAX, 0,
               MPI_COMM_WORLD);
335
336
337     if (rank == 0) {
338         printf("Initial_Local_Serial_Sort_(Max):_%.f_seconds\n", max_serial_time);
339         printf("distribution_time_(Max):_%.f_seconds.\n", max_distr_time);
340         printf("Parallel_Quicksort_Phase_(Max):_%.f_seconds.\n", max_global_sort_time);
341         printf("gather_on_root_time_(Max):_%.f_seconds.\n", max_gather_time);
342         printf("Total_Execution_Time_(Max):_%.f_seconds.\n", max_overall_time);
343
344         check_and_print(all_elements, total_n, output_name);
345         free(all_elements);
346     }
347
348
349     MPI_Finalize();
350     return 0;
351 }
352
353 pivot.c:
354 #include "pivot.h"
355 #include <stdlib.h>
356 #include <string.h>
357 #include <stdio.h>
358 #include <mpi.h>
359
360 int compare(const void* v1, const void* v2){
361     return (*(int*)v1 - *(int*)v2);
362 }
363
364 int get_median(int* elements, int n) {
365     if (n == 0) return 0;
366     if (n % 2 == 0) {

```

```

366         return elements[n / 2 - 1];
367     } else {
368         return elements[n / 2];
369     }
370 }
371
372 int get_larger_index(int *elements, int n, int val) {
373     for (int i = 0; i < n; i++) {
374         if (elements[i] > val) return i;
375     }
376     return n;
377 }
378
379 int select_pivot_median_root(int *elements, int n, MPI_Comm
comm, int *pivot_value) {
380     int rank;
381     MPI_Comm_rank(comm, &rank);
382     int pivot_val = 0;
383
384     if (rank == 0) {
385         if (n > 0) {
386             pivot_val = get_median(elements, n);
387         }
388     }
389     MPI_Bcast(&pivot_val, 1, MPI_INT, 0, comm);
390
391     *pivot_value = pivot_val;
392     return get_larger_index(elements, n, pivot_val);
393 }
394
395 int select_pivot_mean_median(int *elements, int n, MPI_Comm
comm, int *pivot_value) {
396     int rank, size;
397     MPI_Comm_rank(comm, &rank);
398     MPI_Comm_size(comm, &size);
399
400     int local_median = 0;
401     int has_elements = (n > 0) ? 1 : 0;
402     if (n > 0) {
403         local_median = get_median(elements, n);
404     }
405
406     int* all_medians = NULL;
407     int* has_elements_array = NULL;
408     if (rank == 0) {
409         all_medians = (int*)malloc(size * sizeof(int));
410         has_elements_array = (int*)malloc(size * sizeof(int)
);
411         if (!all_medians || !has_elements_array) {

```



```

412         perror("Rank_0: Malloc failed in  

413             select_pivot_mean_median");  

414         MPI_Abort(MPI_COMM_WORLD, 1);  

415     }  

416 }  

417 MPI_Gather(&local_median, 1, MPI_INT, all_medians, 1,  

418           MPI_INT, 0, comm);  

419 MPI_Gather(&has_elements, 1, MPI_INT, has_elements_array  

420           , 1, MPI_INT, 0, comm);  

421  

422 int pivot_val = 0;  

423 if (rank == 0) {  

424     long long sum = 0;  

425     int count = 0;  

426     for (int i = 0; i < size; i++) {  

427         if (has_elements_array[i]) {  

428             sum += all_medians[i];  

429             count++;  

430         }  

431     }  

432     pivot_val = (count > 0) ? (int)(sum / count) : 0;  

433     free(all_medians);  

434     free(has_elements_array);  

435 }  

436 MPI_Bcast(&pivot_val, 1, MPI_INT, 0, comm);  

437  

438 *pivot_value = pivot_val;  

439 return get_larger_index(elements, n, pivot_val);  

440 }  

441  

442 int select_pivot_median_median(int *elements, int n,  

443 MPI_Comm comm, int *pivot_value) {  

444     int rank, size;  

445     MPI_Comm_rank(comm, &rank);  

446     MPI_Comm_size(comm, &size);  

447  

448     int local_median = 0;  

449     int has_elements = (n > 0) ? 1 : 0;  

450     if (n > 0) {  

451         local_median = get_median(elements, n);  

452     }  

453  

454     int* all_medians = NULL;  

455     int* has_elements_array = NULL;  

456     if (rank == 0) {  

457         all_medians = (int*)malloc(size * sizeof(int));  

458         has_elements_array = (int*)malloc(size * sizeof(int));  

459     };  

460     if (!all_medians || !has_elements_array) {

```

```

457         perror("Rank_0: Malloc failed in  

458             select_pivot_median_median");  

459         MPI_Abort(MPI_COMM_WORLD, 1);  

460     }  

461 }  

462 MPI_Gather(&local_median, 1, MPI_INT, all_medians, 1,  

463     MPI_INT, 0, comm);  

464 MPI_Gather(&has_elements, 1, MPI_INT, has_elements_array  

465     , 1, MPI_INT, 0, comm);  

466  

467 int pivot_val = 0;  

468 if (rank == 0) {  

469     int* valid_medians = (int*)malloc(size * sizeof(int)  

470     );  

471     if (!valid_medians) {  

472         perror("Rank_0: Malloc failed for valid_medians"  

473         );  

474         MPI_Abort(MPI_COMM_WORLD, 1);  

475     }  

476     int valid_count = 0;  

477  

478     for (int i = 0; i < size; i++) {  

479         if (has_elements_array[i]) {  

480             valid_medians[valid_count++] = all_medians[i]  

481             ];  

482         }  

483     }  

484  

485     if (valid_count > 0) {  

486         qsort(valid_medians, valid_count, sizeof(int),  

487             compare);  

488         pivot_val = get_median(valid_medians,  

489             valid_count);  

490     }  

491  

492     free(all_medians);  

493     free(has_elements_array);  

494     free(valid_medians);  

495 }  

496 MPI_Bcast(&pivot_val, 1, MPI_INT, 0, comm);  

497  

498 *pivot_value = pivot_val;  

499 return get_larger_index(elements, n, pivot_val);  

500 }  

501  

502 int select_pivot_smallest_root(int *elements, int n,  

503 MPI_Comm comm, int *pivot_value) {  

504     int rank;  

505     MPI_Comm_rank(comm, &rank);

```

```

498     int pivot_val = 0;
499
500     if (rank == 0 && n > 0) {
501         pivot_val = elements[0];
502     }
503     MPI_Bcast(&pivot_val, 1, MPI_INT, 0, comm);
504
505     *pivot_value = pivot_val;
506     return get_larger_index(elements, n, pivot_val);
507 }
508
509 int select_pivot(int pivot_strategy, int *elements, int n,
510 MPI_Comm communicator, int *pivot_value) {
511     int pivot_index_result = 0;
512
513     switch (pivot_strategy) {
514         case MEDIAN_ROOT:
515             pivot_index_result = select_pivot_median_root(
516                 elements, n, communicator, pivot_value);
517             break;
518         case MEAN_MEDIAN:
519             pivot_index_result = select_pivot_mean_median(
520                 elements, n, communicator, pivot_value);
521             break;
522         case MEDIAN_MEDIAN:
523             pivot_index_result = select_pivot_median_median(
524                 elements, n, communicator, pivot_value);
525             break;
526         default: // SMALL_ROOT
527             pivot_index_result = select_pivot_smallest_root(
528                 elements, n, communicator, pivot_value);
529             break;
530     }
531     return pivot_index_result;
532 }

```